



307401  
Big Data and Data Warehouses  
Running MapReduce Using MRJob Python Library



# Implement MapReduce Jobs using MRJOB

Java is the standard language for running MapReduce jobs.

**mrjob** is a Python library developed by Yelp that lets us:

- Write MapReduce jobs entirely in Python.
- **mrjob** can run MapReduce jobs locally, on Hadoop, or on EMR (AWS) with the same code.
- It allows us to focus on logic, not infrastructure.

# MRJob example: Words Count

Python MapReduce code to count the words in our file Input.txt

```
from mrjob.job import MRJob

class MRWordCount(MRJob):
    def mapper(self, _, line):
        for word in line.split():
            yield word.lower(), 1

    def reducer(self, word, counts):
        yield word, sum(counts)

if __name__ == '__main__':
    MRWordCount.run()
```

```
input.txt file
hello world bye world
hello hadoop mapreduce
world
```

Part	Description
from mrjob.job import MRJob	Imports the MRJob base class used to define MapReduce jobs.
class MRWordCount(MRJob):	Defines a new job called MRWordCount that inherits from MRJob.
def mapper(self, _, line):	The <b>mapper</b> function runs on each line of input text.
for word in line.split(): yield word, 1	Splits the line into words and emits each word paired with the number 1.
def reducer(self, word, counts):	The <b>reducer</b> function takes all counts for each word.
yield word, sum(counts)	Adds up all counts for a word and emits the total.

# Step-by-Step: Who Does What

**lifecycle of a key–value pair:**

- **Mapper emits (key, value):**

- This happens in the *mapper process* (a JVM container on a node).
- These pairs are first kept in memory in a buffer.
- When the buffer fills, Hadoop **spills** the data to **local disk** in sorted order by key.

- **Mapper output → local disk:**

- Each mapper writes one or more **spill files**.
- When the mapper finishes, it **merges** all spill files into one final sorted output file (still on the mapper's local disk).

- **Shuffle and transfer phase:**

- The **reducer processes** contact each mapper node (over HTTP) to fetch the relevant partitions of the mapper outputs.
- This is managed by the **Hadoop framework**, not your code.
- Data is transferred over the network: each reducer fetches the partitions of intermediate data corresponding to its assigned key range.

- **Sort and merge on the reducer side:**

- Each reducer merges all fetched mapper partitions, sorts by key again (if needed), and groups values for each key.
- Then it calls your reducer function for each group.

- **Reducer output:**

- The reducer writes final results to **HDFS** (Hadoop Distributed File System).

## Run Code Locally

1) Download and install Python version **3.10 or 3.11** (they support MRJob) – Look for [Windows installer \(64-bit\)](#)  
<https://www.python.org/downloads/release/python-3119/>

2) Open the command line terminal and install mrjob on your local computer:

```
pip install mrjob
```

3) Download words\_count.py and input.txt files from BlackBoard.

4) Open the command line terminal and type:

```
python words_count.py input.txt
```

You should see:

```
"bye"      1
"hadoop"   1
"hello"    2
"mapreduce" 1
"world"    3
```

This verifies that the logic works before distributing the job.

## Understanding What Happened

- Hadoop divided your file into chunks (input splits).
- Each **mapper** processed one split, generating intermediate (word, 1) pairs.
- Hadoop grouped identical keys (all “hello”s together).
- The **reducer** summed values for each key.
- This is the essence of **MapReduce** — distributed computation through key-value pair processing.

# Movie Ratings Count Example

User ID	Movie ID	Rating	Time Stamp
0	50	5	881250949
0	172	5	881250949
0	133	1	881250949
196	242	3	881250949
186	302	3	891717742
22	377	1	878887116
244	51	2	880606923
166	346	1	886397596
298	474	4	884182806
115	265	2	881171488
253	465	5	891628467
305	451	3	886324817





# Using MRStep (for multi-step jobs)

When you have multiple map/reduce phases, use MRStep, for example, counting word frequency then sorting by frequency.

```
from mrjob.job import MRJob
from mrjob.step import MRStep

class RatingsBreakdown(MRJob):
    def steps(self):
        return [MRStep(mapper=self.mapper_get_ratings,
                       reducer=self.reducer_count_ratings),
                MRStep(reducer=self.reducer_sort_by_count)]

    def mapper_get_ratings(self, _, line):
        # Input: userID, movieID, rating, timestamp (tab-separated)
        try:
            userID, movieID, rating, timestamp = line.strip().split('\t')
            yield rating, 1
        except ValueError:
            pass # skip malformed lines

    def reducer_count_ratings(self, rating, counts):
        # Count how many times each rating appears
        yield None, (sum(counts), rating)

    def reducer_sort_by_count(self, _, count_rating_pairs):
        # Sort by count (descending)
        sorted_pairs = sorted(count_rating_pairs, reverse=True)
        for count, rating in sorted_pairs:
            yield rating, count

if __name__ == '__main__':
    RatingsBreakdown.run()
```

## MRStep in mrjob:

- MRStep defines one MapReduce phase consisting of an optional mapper, combiner, and reducer.
- It allows chaining multiple MapReduce steps within a single MRJob class.
- The steps() method returns a list of MRStep objects executed sequentially.
- Output from one step automatically becomes input for the next step.
- **Each MRStep runs as a separate MapReduce job under the hood.**
- **Common use: perform an initial computation (e.g., counting) in the first step, then process or sort those results in a later step.**

In the example:

- Step 1 mapper emits (rating, 1); reducer counts occurrences and emits (None, (count, rating)).
- Step 2 reducer receives all (count, rating) pairs, sorts them, and emits (rating, count).
- Using MRStep avoids writing multiple separate jobs and enables building pipelines inside one mrjob script.

Step	Function	Input Example	Operation / Logic	Output Example
<b>1. Mapper</b>	<code>mapper_get_ratings(_, line)</code>	Each line of the file: 1 31 2.5 1260759144	Splits each line by tab into userID, movieID, rating, timestamp. Emits (rating, 1) for each line.	(2.5, 1)
<b>2. Shuffle &amp; Sort (automatic)</b>	<i>(Hadoop/MRJob built-in)</i>	Mapper outputs: (2.5,1), (3.0,1), (2.5,1)	Groups values by key (rating) and sorts keys in ascending order before passing to reducers.	2.5 → [1,1], 3.0 → [1]
<b>3. Reducer #1</b>	<code>reducer_count_ratings(rating, counts)</code>	rating = "2.5", counts = [1,1]	Sums the counts for each rating. Emits all results under a single key None so they go to the same reducer in the next step.	(None, (2, "2.5"))
<b>4. Shuffle #2 (automatic)</b>	<i>(Hadoop/MRJob built-in)</i>	All pairs (None, (count, rating))	Since all share the same key None, they are sent together to one reducer.	[ (2, "2.5"), (1, "3.0"), ... ]
<b>5. Reducer #2</b>	<code>reducer_sort_by_count(_, count_rating_pairs)</code>	[(2, "2.5"), (1, "3.0"), (5, "4.0")]	Sorts all pairs by count in descending order. Emits (rating, count) for each sorted pair.	"4.0" → 5, "2.5" → 2, "3.0" → 1
<b>6. Final Output</b>	—	—	Displays ratings sorted by how often they appear.	...
4.0 5				
2.5 2				
3.0 1				



# Running MapReduce on Hadoop

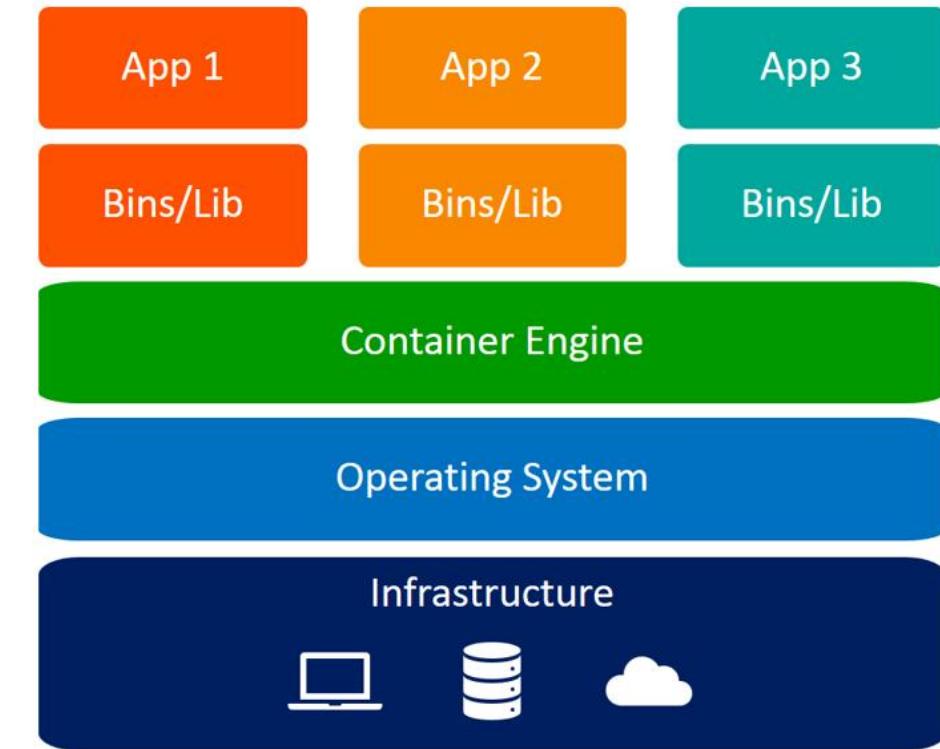
# Introduction to Docker and Docker Compose

## What is Docker

- Docker is a containerization platform that allows you to package an application together with all its dependencies into a single portable unit called a container.
- Containers ensure that the application runs the same way on any system, regardless of the underlying operating system or configuration.
- Unlike virtual machines, Docker containers share the host system's operating system kernel, making them lightweight, efficient, and fast to start.
- Analogy: A Docker container is like a sealed lab box that includes everything your program needs to run.

## Why We Use Docker in This Lab

- To create a ready-to-use Hadoop environment without complex installation steps
- To prevent version or dependency conflicts between systems
- To make the lab easily reproducible on student laptops or classroom machines
- To simulate a cluster setup using a single machine



Containers

# Installing and Configuring Docker

## Step 1: Install Docker Desktop

1. Go to <https://www.docker.com/products/docker-desktop>
2. Download and install **Docker Desktop for Windows**.
3. During setup, **enable the WSL 2 backend**.
4. Restart your computer.
5. Verify Docker is running — the **whale icon** should appear in the system tray.

## Step 2: Adjust Docker Resources

1. Open Docker Desktop → Settings → Resources.
2. Allocate:
  - CPUs: 2
  - Memory: 4–6 GB
  - Swap: 1–2 GB
3. Click **Apply & Restart**.

Proper resource allocation is critical. Hadoop requires enough memory for its daemons (NameNode, DataNode, ResourceManager, etc.) to start successfully.

# 1) Download and Run the Hadoop Container

Open the command line and the following commands:

a. Pull Hadoop container from the Internet ([Docker Hub](#)) - open a Windows command line terminal and type the command below:

```
docker pull msfasha/hadoop-petra:latest
```

b. Run the container on you machine - open a Windows command line terminal and type the command below:

```
docker run -it -p 9870:9870 -p 8088:8088 --name hadoop-lab msfasha/hadoop-petra
```

c. We can check the running Hadoop environment using Web UI:

HDFS UI: <http://localhost:9870>

YARN UI: <http://localhost:8088>

d. Enter the container shell - open a Windows command line terminal and type the command below:

```
docker exec -it hadoop-lab bash
```

Basic test commands you can run inside Hadoop container:

```
# check Hadoop daemons
```

```
jps
```

```
# list files in HDFS
```

```
hdfs dfs -ls /
```

```
# check running containers or YARN jobs
```

```
yarn application -list
```



## 2) Copy Code File and Data File Into the Container

a. Download (**words\_count.py** and **inputs.txt**) from Black Board.

b. Copy files from your machine into the docker container - open a Windows command line terminal and type the command below:

```
docker cp words_count.py hadoop-lab:/opt/hadoop/
```

```
docker cp input.txt hadoop-lab:/opt/hadoop/
```

c. Enter the docker container – type the command below inside Windows command line terminal:

```
docker exec -it hadoop-lab bash
```

c. Copy data file from the container into Hadoop HDFS – type the command below in Hadoop container command line:

```
hdfs dfs -mkdir /input
```

```
hdfs dfs -put /opt/hadoop/input.txt /input/
```

### 3) Run the MapReduce Job in Hadoop

#### a. Run the Python MapReduce Program on Hadoop -

```
python3 /opt/hadoop/words_count.py -r hadoop hdfs://input/input.txt -o hdfs://output
```

Explanation:

- -r hadoop: Run using Hadoop's MapReduce engine.
- hdfs://input/input.txt: Input file in HDFS.
- -o hdfs://output\_mrjob: Output folder in HDFS.
- Monitor progress in YARN's web interface: <http://localhost:8088>

## 4) Check the Results

- a. When the job completes, check the output in HDFS:

```
hdfs dfs -cat /output/part-00000
```

Expected output:

```
"bye" 1
"hadoop" 1
"hello" 2
"mapreduce" 1
"world" 3
```

## Key Takeaways

- Hadoop allows distributed processing using MapReduce.
- Docker provides an easy way to run Hadoop without complex installation.
- HDFS stores data across multiple nodes; YARN manages job execution.
- mrjob makes writing Python MapReduce programs straightforward.
- Hadoop web interfaces offer valuable visualization and debugging tools.